

ML-Classification-Categorical-Snow

November 30, 2017

```
In [1]: from pyspark.sql.types import StructType, StructField, FloatType, LongType, StringType
        from pyspark.shell import spark
```

```
feats = []
f = open('features.txt')
for line_num, line in enumerate(f):
    if line_num == 0:
        # Timestamp
        feats.append(StructField(line.strip(), LongType(), True))
    elif line_num == 1:
        # Geohash
        feats.append(StructField(line.strip(), StringType(), True))
    else:
        # Other features
        feats.append(StructField(line.strip(), FloatType(), True))

schema = StructType(feats)
```

Welcome to

```
  _--_
 /  _/  _--_  _--_  _--_  _--_
 \  \ /  _ \  _ \  _ \  _ \  _ \
/_  /  . _/ \_ , _/ /  /_ / \_ \   version 2.2.0
/_  /
```

Using Python version 3.6.3 (default, Oct 6 2017 12:04:38)
SparkSession available as 'spark'.

```
In [2]: df = spark.read.format('csv').option('sep', '\t').schema(schema).load('inputs/mini-samp
```

```
In [5]: col_names = []
        for i in range(2, len(df.columns)):
            col_names.append(df.columns[i])
```

```
In [39]: from pyspark.ml.stat import Correlation
         from pyspark.ml.feature import VectorAssembler
         vectorAssembler = VectorAssembler(inputCols=col_names,
```

```

                                outputCol="features")
trans_features = vectorAssembler.transform(df)
coeff = Correlation.corr(trans_features, 'features', method='pearson').collect()[0][0]

In [43]: mtrx = coeff.toArray()

In [44]: list_corr_pairs_coefs = []
        feature_pairs = []

        for i in range(0,56):
            for j in range(0,56):

                if (i != j) and not(((col_names[i]+"_"+col_names[j]) in feature_pairs) or
                                     ((col_names[j]+"_"+col_names[i]) in feature_pairs)):
                    feature_pairs.append(col_names[i] + "_" + col_names[j])
                    corr_pair_coeff = []
                    corr_pair_coeff.append(col_names[i] + " , " + col_names[j])
                    corr_pair_coeff.append(col_names[i])
                    corr_pair_coeff.append(col_names[j])
                    corr_pair_coeff.append(float(mtrx[i][j]))
                    list_corr_pairs_coefs.append(tuple(corr_pair_coeff))

In [45]: df_corr_coeff_col_names = []
        df_corr_coeff_col_names.append(StructField("Feature_Pair", StringType(), True))
        df_corr_coeff_col_names.append(StructField("Feature1", StringType(), True))
        df_corr_coeff_col_names.append(StructField("Feature2", StringType(), True))
        df_corr_coeff_col_names.append(StructField("Pearson_Coeff", FloatType(), True))
        df_corr_coeff = spark.createDataFrame(list_corr_pairs_coefs, StructType(df_corr_coeff_col_names))

In [46]: df_corr_coeff.where((df_corr_coeff.Feature1 == 'temperature_surface' ) |
                             (df_corr_coeff.Feature2 == 'temperature_surface')).sort(df_corr_coeff.Feature1, df_corr_coeff.Feature2)

```

| Feature_Pair | Feature1 | Feature2 | Pearson_Coeff |
|---|----------------------|----------|---------------|
| upward_long_wave... upward_long_wave... | temperature_surface | | 0.9507364 |
| geopotential_heig... geopotential_heig... | temperature_surface | | 0.88414276 |
| downward_long_wav... downward_long_wav... | temperature_surface | | 0.83087575 |
| geopotential_heig... geopotential_heig... | temperature_surface | | 0.8231031 |
| temperature_surfa... temperature_surface | precipitable_wate... | | 0.708688 |
| temperature_surfa... temperature_surface | convective_availa... | | 0.4671096 |
| temperature_surfa... temperature_surface | geopotential_heig... | | 0.37617406 |
| temperature_surfa... temperature_surface | downward_short_wa... | | 0.34058186 |
| visibility_surfac... visibility_surface | temperature_surface | | 0.28188124 |
| pressure_surface ... pressure_surface | temperature_surface | | 0.20792589 |
| geopotential_heig... geopotential_heig... | temperature_surface | | 0.117955595 |
| lightning_surface... lightning_surface | temperature_surface | | 0.11171691 |
| vegetation_surfac... vegetation_surface | temperature_surface | | 0.09736003 |
| temperature_surfa... temperature_surface | v-component_of_wi... | | 0.08873683 |

```
|upward_short_wave...|upward_short_wave...| temperature_surface| 0.08651417|
|v-component_of_wi...|v-component_of_wi...| temperature_surface| 0.07520971|
|categorical_rain...|categorical_rain...| temperature_surface| 0.05293761|
|temperature_surfa...| temperature_surface|latent_heat_net_f...| 0.012255301|
|sensible_heat_net...|sensible_heat_net...| temperature_surface| 0.011762388|
|drag_coefficient...|drag_coefficient...| temperature_surface| 0.00251122|
+-----+-----+-----+-----+
only showing top 20 rows
```

In [47]: `#categorical_snow_yes1_no0_surface`

```
df_corr_coeff.where((df_corr_coeff.Feature1 == 'categorical_snow_yes1_no0_surface') |
                    (df_corr_coeff.Feature2 == 'categorical_snow_yes1_no0_surface')
                    ).sort(df_corr_coeff.Pearson_Coeff.desc()).show()
```

```
+-----+-----+-----+-----+
|      Feature_Pair|      Feature1|      Feature2|Pearson_Coeff|
+-----+-----+-----+-----+
|categorical_snow...|categorical_snow...| snow_cover_surface| 0.30012408|
|categorical_snow...|categorical_snow...|maximumcomposite...| 0.2645679|
|categorical_snow...|categorical_snow...| snow_depth_surface| 0.2408605|
|categorical_snow...|categorical_snow...|      albedo_surface| 0.2254173|
|water_equiv_of_ac...|water_equiv_of_ac...|categorical_snow...| 0.2251549|
|categorical_snow...|categorical_snow...| pressure_tropopause| 0.2167491|
|relative_humidity...|relative_humidity...|categorical_snow...| 0.18254393|
|categorical_snow...|categorical_snow...|friction_velocity...| 0.17027378|
|categorical_snow...|categorical_snow...|temperature_tropo...| 0.16374761|
|categorical_snow...|categorical_snow...|plant_canopy_surf...| 0.124861635|
|number_of_soil_la...|number_of_soil_la...|categorical_snow...| 0.11009545|
|categorical_snow...|categorical_snow...|land_cover_land1...| 0.103146575|
|categorical_snow...|categorical_snow...|soil_porosity_sur...| 0.103146575|
|categorical_snow...|categorical_snow...|surface_wind_gust...| 0.098745205|
|categorical_snow...|categorical_snow...|transpiration_str...| 0.09675234|
|categorical_snow...|categorical_snow...|geopotential_heig...| 0.094172865|
|categorical_snow...|categorical_snow...|surface_roughness...| 0.08816603|
|categorical_snow...|categorical_snow...|geopotential_heig...| 0.07994601|
|categorical_snow...|categorical_snow...|v-component_of_wi...| 0.07057235|
|categorical_snow...|categorical_snow...|soil_type_as_in_z...| 0.06916345|
+-----+-----+-----+-----+
only showing top 20 rows
```

In [48]: `df_corr_coeff.where((df_corr_coeff.Feature1 == 'categorical_snow_yes1_no0_surface') |
 (df_corr_coeff.Feature2 == 'categorical_snow_yes1_no0_surface')
).sort(df_corr_coeff.Pearson_Coeff.asc()).show()`

```
+-----+-----+-----+-----+
|      Feature_Pair|      Feature1|      Feature2|Pearson_Coeff|
```

```

+-----+-----+-----+-----+
|visibility_surfac...| visibility_surface|categorical_snow...| -0.45647138|
|categorical_snow...|categorical_snow...|geopotential_heig...| -0.26098943|
|categorical_snow...|categorical_snow...| temperature_surface| -0.21268843|
|categorical_snow...|categorical_snow...|upward_long_wave...| -0.2037285|
|geopotential_heig...|geopotential_heig...|categorical_snow...| -0.19996537|
|categorical_snow...|categorical_snow...|precipitable_wate...| -0.15940763|
|pressure_reduced...|pressure_reduced...|categorical_snow...| -0.13463028|
|categorical_snow...|categorical_snow...|mean_sea_level_pr...| -0.122622825|
|categorical_snow...|categorical_snow...|downward_long_wav...| -0.11768969|
|categorical_snow...|categorical_snow...| pressure_surface| -0.10548218|
|categorical_snow...|categorical_snow...|convective_availa...| -0.07774716|
|categorical_snow...|categorical_snow...|downward_short_wa...| -0.07093287|
|categorical_snow...|categorical_snow...| vegetation_surface| -0.05607693|
|categorical_snow...|categorical_snow...|categorical_rain...| -0.052368812|
|categorical_snow...|categorical_snow...|v-component_of_wi...| -0.019803138|
|categorical_snow...|categorical_snow...|geopotential_heig...| -0.011431282|
|categorical_snow...|categorical_snow...|latent_heat_net_f...| -0.00841602|
|sensible_heat_net...|sensible_heat_net...|categorical_snow...| -0.008250166|
|categorical_freez...|categorical_freez...|categorical_snow...| -0.004402432|
|upward_short_wave...|upward_short_wave...|categorical_snow...| -0.003501976|
+-----+-----+-----+-----+
only showing top 20 rows

```

```

In [49]: #snow_cover_surface
         #snow_depth_surface
         #visibility_surface

```

```

df_corr_coeff.where(((df_corr_coeff.Feature1 == 'categorical_snow_yes1_no0_surface') &
                      (df_corr_coeff.Feature2 == 'temperature_surface')) |
                    ((df_corr_coeff.Feature2 == 'categorical_snow_yes1_no0_surface') &
                      (df_corr_coeff.Feature1 == 'temperature_surface')))
                ).sort(df_corr_coeff.Pearson_Coeff.asc()).show()

#temperature_surface

```

```

+-----+-----+-----+-----+
|      Feature_Pair|      Feature1|      Feature2|Pearson_Coeff|
+-----+-----+-----+-----+
|categorical_snow...|categorical_snow...|temperature_surface| -0.21268843|
+-----+-----+-----+-----+

```

```

In [50]: df_corr_coeff.where(((df_corr_coeff.Feature1 == 'categorical_snow_yes1_no0_surface') &
                               (df_corr_coeff.Feature2 == 'precipitable_water_entire_atmosphere') &
                               ((df_corr_coeff.Feature2 == 'categorical_snow_yes1_no0_surface') &

```

```

(df_corr_coeff.Feature1 == 'precipitable_water_entire_atmosphere'
).sort(df_corr_coeff.Pearson_Coeff.asc()).show()
#precipitable_water_entire_atmosphere

```

| Feature_Pair | Feature1 | Feature2 | Pearson_Coeff |
|----------------------|----------------------|----------------------|---------------|
| categorical_snow_... | categorical_snow_... | precipitable_wate... | -0.15940763 |

```

In [51]: #snow_cover_surface
#snow_depth_surface
#visibility_surface
#temperature_surface
#precipitable_water_entire_atmosphere

```

```

In [61]: def prepare_data(dframe, predictors, target):
    assembler = VectorAssembler(inputCols=predictors, outputCol="features")
    output = assembler.transform(dframe)
    return output.select("features", target).withColumnRenamed(target, "label")

```

```

prepped = prepare_data(df,
    ["snow_cover_surface",
     "snow_depth_surface",
     "visibility_surface",
     "temperature_surface",
     "precipitable_water_entire_atmosphere"],
    "categorical_snow_yes1_no0_surface")

```

```

prepped.show()
(trainingData, testData) = prepped.randomSplit([0.9, 0.1])

```

| features | label |
|-----------------------|-------|
| [0.0,0.0,24221.20...] | 0.0 |
| [100.0,0.01775999...] | 0.0 |
| [100.0,0.05375999...] | 0.0 |
| [0.0,0.0,22024.20...] | 0.0 |
| [0.0,0.0,24222.64...] | 0.0 |
| [100.0,0.90279996...] | 0.0 |
| [0.0,0.0,24225.47...] | 0.0 |
| [0.0,0.0,24225.27...] | 0.0 |
| [0.0,0.0,24224.25...] | 0.0 |
| [0.0,0.0,24225.79...] | 0.0 |
| [0.0,0.0,21825.27...] | 0.0 |

```
| [0.0,0.0,24222.57...| 0.0|
| [100.0,0.00944000...| 0.0|
| [0.0,0.0,24225.83...| 0.0|
| [100.0,0.02439999...| 0.0|
| [0.0,0.0,24224.25...| 0.0|
| [0.0,0.0,24222.07...| 0.0|
| [0.0,0.0,24224.07...| 0.0|
| [0.0,0.0,24222.07...| 0.0|
| [0.0,0.0,24223.87...| 0.0|
```

```
+-----+
```

only showing top 20 rows

```
In [67]: from pyspark.ml.classification import RandomForestClassifier
         from pyspark.ml.evaluation import BinaryClassificationEvaluator

         rf = RandomForestClassifier(numTrees=3, maxDepth=2, labelCol="label", seed=42)
         model = rf.fit(prepped)
```

```
In [68]: model.featureImportances
```

```
Out[68]: SparseVector(5, {0: 0.075, 1: 0.0754, 2: 0.7194, 3: 0.1303})
```

```
In [72]: predictions = model.transform(testData)
         predictions.select(predictions.probability).show(n=5)
```

```
+-----+
|      probability|
+-----+
| [0.95664279808010...|
| [0.75998844918654...|
| [0.95664279808010...|
| [0.95664279808010...|
| [0.95664279808010...|
+-----+
```

only showing top 5 rows

```
In [74]: predictions.select(*predictions.columns).show(n=5)
```

```
+-----+-----+-----+-----+-----+
|      features|label|rawPrediction|probability|prediction|
+-----+-----+-----+-----+-----+
| [0.0,0.0,17.33257...| 0.0|[2.86992839424032...|[0.95664279808010...| 0.0|
| [0.0,0.0,17.51190...| 0.0|[2.27996534755964...|[0.75998844918654...| 0.0|
| [0.0,0.0,17.51190...| 0.0|[2.86992839424032...|[0.95664279808010...| 0.0|
| [0.0,0.0,18.01349...| 0.0|[2.86992839424032...|[0.95664279808010...| 0.0|
```

```
|[0.0,0.0,18.04895...| 0.0|[2.86992839424032...|[0.95664279808010...| 0.0|
+-----+-----+-----+-----+
only showing top 5 rows
```

```
In [84]: evaluator = BinaryClassificationEvaluator()
         evaluator.evaluate(predictions)#by default considers area under roc to evaluate
```

```
Out[84]: 0.965504638264051
```

```
In [87]: samp_predns = predictions.sample(False, .1)
         samp_predns
```

```
Out[87]: DataFrame[features: vector, label: float, rawPrediction: vector, probability: vector,
```

```
In [97]: import numpy as np
         from sklearn import metrics
         predns = np.array(samp_predns.select(samp_predns.prediction).collect())
```

```
In [98]: labels = np.array(samp_predns.select(samp_predns.label).collect())
```

```
In [99]: fpr, tpr, thresholds = metrics.roc_curve(labels, predns)
         roc_auc = metrics.auc(fpr, tpr)
```

```
In [100]: import matplotlib.pyplot as plt
          plt.title('Area under ROC Curve')
          plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
          plt.legend(loc = 'lower right')
          plt.plot([0, 1], [0, 1], 'r--')
          plt.xlim([0, 1])
          plt.ylim([0, 1])
          plt.ylabel('True Positive Rate')
          plt.xlabel('False Positive Rate')
          plt.show()
```

